

# Final Review

# Database APIs

Database libraries: Provide an API to a DB

Cursors: Iteratively page through results

“Impedance mismatches”

Types: different ints, floats, dates, etc

Object: Objects/structs != rows

Constraints: program logic VS DB rules

SQL Injection: Building strings instead of letting API do it

# Duplication is bad

One type of meat, cheese, and vegetable

Pizza	Topping	Type
1	Mozzarella	Cheese
1	Pepperoni	Meat
1	Olives	Vegetable
2	Mozzarella	Cheese
2	Sausage	Meat
2	Peppers	Vegetable

Key? (Pizza, Type)

# Anomalies

Update: Correct a type, must update all rows

Insert: Add a type without a pizza?

Delete: Cancel a pizza; lose a type!

# Decomposition

Replace schema R with 2+ smaller schemas that

1. each contain subset of attrs in R
2. together include all attrs in R

ABCD replaced with AB, BCD or AB, BC, CD

## Desirable properties

1. Lossless join: able to recover R from smaller relations
2. Dependency preserving: enforce constraints on R by only enforcing constraints on smaller schemas (no joins)

# Functional Dependencies (FD)

$$X \rightarrow Y$$

holds on R

if  $t_1.X = t_2.X$  then  $t_1.Y = t_2.Y$

where  $X, Y$  are subsets of attrs in R

## Examples of FDs in person-hobbies table

$sid \rightarrow name, address$

$hobby \rightarrow cost$

$sid, hobby \rightarrow name, address cost$

# Keys and Functional Dependencies

A key must determine all column values

Must be the left hand side (aka input or source) of FDs

$BC \rightarrow N, N \rightarrow BO$

NBO key is N: determines all other fields

Banker name determines branch, office

NBO is in BCNF: All FDs are keys

# Where do FDs come from?

thinking really hard aka application semantics  
can't stare at database to derive (like ICs)

Like a Mathematics conjecture:

one counter example can disprove, but examples can't prove  
there are no example in the universe



# Closure of FDs

If I know

$\text{Name} \rightarrow \text{Bday}$  and  $\text{Bday} \rightarrow \text{age}$

Then it implies

$\text{Name} \rightarrow \text{age}$

An FD  $f'$  is implied by set  $F$  if  $f'$  is true when  $F$  is true

$F^+$ : the **closure** of  $F$  is all FDs implied by  $F$

Can we construct this closure automatically? YES

# Closure of FDs

*Inference rules* called **Armstrong's Axioms**

Reflexivity      if  $Y \subseteq X$  then  $X \rightarrow Y$

Augmentation    if  $X \rightarrow Y$  then  $XZ \rightarrow YZ$  for any  $Z$

Transitivity      if  $X \rightarrow Y$  &  $Y \rightarrow Z$  then  $X \rightarrow Z$

These are **sound** and **complete** rules

sound            doesn't produce FDs not in the closure

complete        doesn't miss any FDs in the closure

# Closure of FDs

$F = \{A \rightarrow B, B \rightarrow C, CB \rightarrow E\}$

Is  $A \rightarrow E$  in the closure?

$A \rightarrow B$

given

$A \rightarrow AB$

augmentation  $A$

$A \rightarrow BB$

apply  $A \rightarrow B$

$A \rightarrow BC$

apply  $B \rightarrow C$

$BC \rightarrow E$

given

$A \rightarrow E$

transitivity

# Minimum Cover of FDs

## 1. Turn FDs into *standard form*

Split FDs so there is one attribute on the right side

## 2. Minimize left side of each FD

For each FD, check if can delete a left attribute using another FD  
given  $ABC \rightarrow D, B \rightarrow C$  can reduce to  $AB \rightarrow D, B \rightarrow C$

## 3. Delete redundant FDs

check each remaining FD and see if it can be deleted  
e.g., in closure of the other FDs

**2 must happen before 3!**

# Normal Forms

Criteria met by a relation  $R$  wrt functional dependencies

Boyce Codd Normal Form (BCNF)

No redundancy, may lose dependencies

Third Normal Form (3NF)

May have redundancy, no decomposition problems

# Lossless Join Decomposition

join the decomposed tables to get *exactly the* original

e.g., decompose  $R$  into tables  $X, Y$

$$\pi_X(R) \bowtie \pi_Y(R) = R$$

Lossless wrt  $F$  if and only if  $F^+$  contains

$$X \cap Y \rightarrow X \text{ or } Y \cap X \rightarrow Y$$

intersection of  $X, Y$  is a key for one of them

# Lossless Join Decomposition

Lossless wrt  $F$  if and only if  $F^+$  contains

$$X \cap Y \rightarrow X \text{ or } Y \cap X \rightarrow Y$$

intersection of  $X, Y$  is a key for one of them

FDs:  $A \rightarrow C, A \rightarrow B$

A	B	C
1	2	1
5	3	4
9	2	6



A	B
1	2
5	3
9	2

B	C
2	1
3	4
2	6



A	B	C
1	2	1
5	3	4
9	2	6
1	2	6
9	2	1

Lossy!  $AB \cap BC = B$  doesn't determine anything

# Dependency-preserving Decomposition

$F_R$  = Projection of  $F$  onto  $R$

FDs  $X \rightarrow Y$  in  $F^+$  s.t.  $X$  and  $Y$  attrs are in  $R$

Subset of  $F$  that are “valid” for  $R$

If  $R$  decompose to  $X, Y$ .

FDs that hold on  $X, Y$  equivalent to all FDs on  $R$

$$(F_X \cup F_Y)^+ = F^+$$

Consider  $ABCD$ ,  $C$  is key,  $AB \rightarrow C$ ,  $D \rightarrow A$

BCNF decomposition:  $BCD, DA$

$AB \rightarrow C$  doesn't apply to either table!



# BCNF

Relation R in BCNF has *no redundancy* wrt FDs  
(only FDs are key constraints)

F: set of functional dependencies over relation R

X: Subset of attributes of R

A: One attribute of R

for  $(X \rightarrow A)$  in F

A is in X (trivial/reflexivity) OR

X is a superkey of R

(superkeys include candidate keys)

# BCNF

while BCNF is violated

  R with FDs  $F_R$

  if  $X \rightarrow Y$  violates BCNF

    turn R into R-Y & XY

# 3NF

$F^{\min}$  = minimal cover of  $F$

Run BCNF using  $F^{\min}$

for  $X \rightarrow Y$  in  $F^{\min}$  not in projection onto  $R_1 \dots R_N$

create relation  $XY$

BCNO       $BC \rightarrow N, N \rightarrow BO$

NBO, CN using  $N \rightarrow BO$

# Disks

Time to access (read or write) a disk block

seek time	2-4 msec avg (arm movement)
rotational delay	2-4 msec (based on rotation speed)
transfer time	0.3 msec/64kb page

Throughput

read	~150 MB/sec
write	~50 MB/sec

Key: reduce seek and rotational delays  
HW & SW approaches

# What is a page?

Unit of transfer between storage and database

Typically fixed size

Small enough for one I/O to be fast

Big enough to not be wasteful

Usually a multiple of 4 kB

- Intel virtual memory hardware page size

- Modern disk sector size (minimum I/O size)

# Records and Files

Record: “application” storage unit

e.g. a row in a table

Page: Collection of records

File: Collection of pages

insert/delete/modify record

get(record\_id) a record

scan all records

May be in multiple OS files spanning multiple disks

# Unordered Heap Files

Unordered collection of records

Pages allocated as collection grows

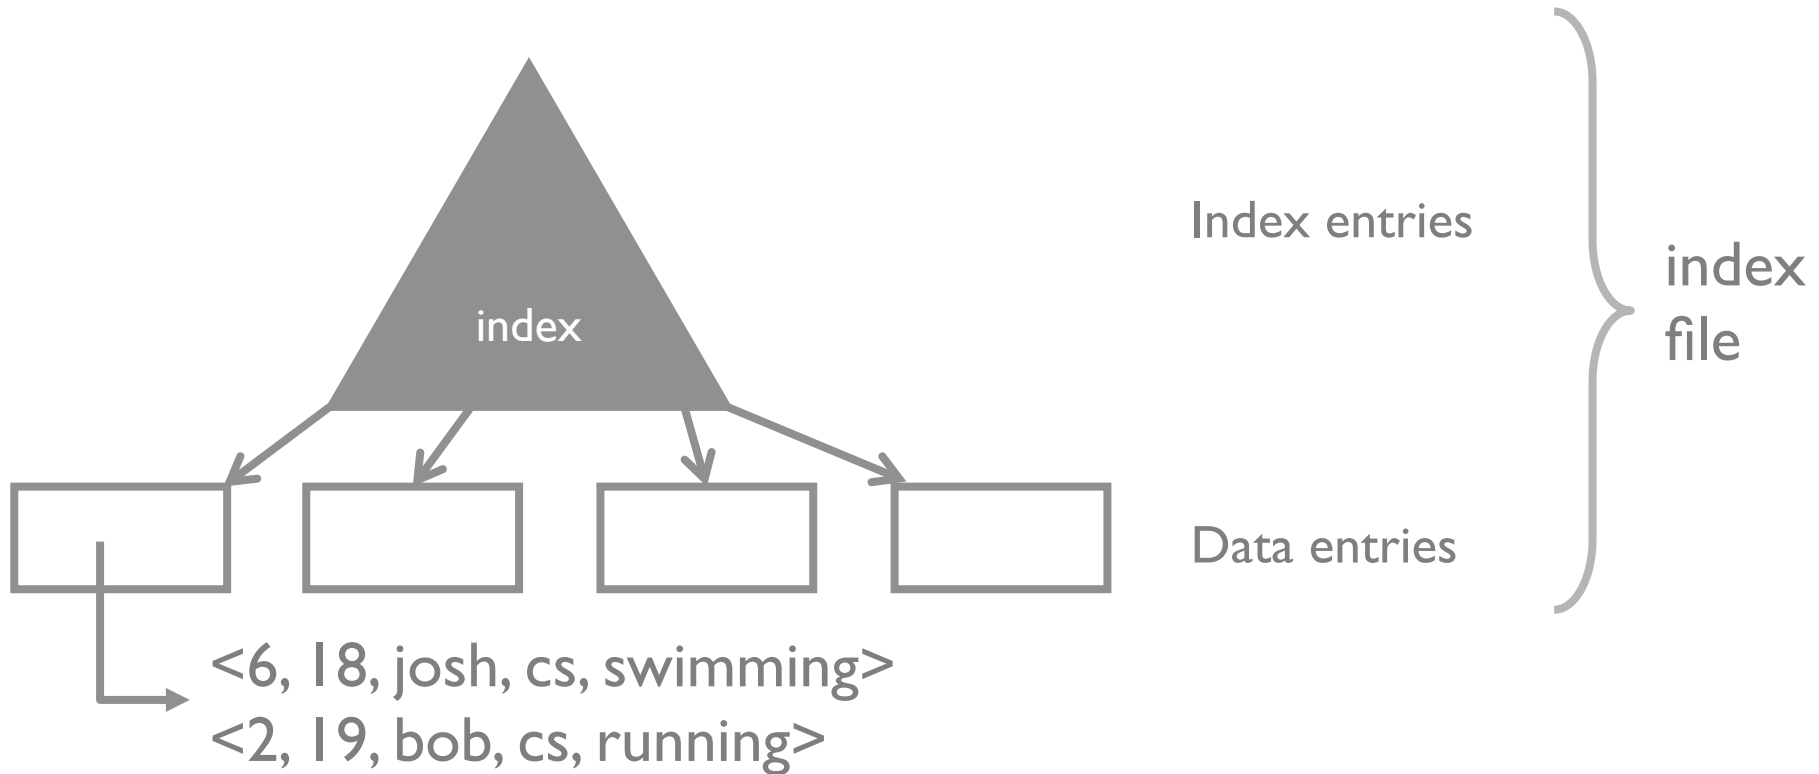
Need to track:

- pages in file

- free space on pages

- records on page


# High level (Primary) index structure



What is a data entry?  
actual data record




# High level (Secondary) index structure



What is a data entry?  
actual data record  
<search key value, rid>

Tradeoffs  
directly access tuple.  
compact, fixed size entries

# B+ Tree on (age, name)




How do the following queries use the index on (age, name)?

SELECT age WHERE age = 14

SELECT \* WHERE age < 18 AND name < 'monica'

SELECT age WHERE name = 'bobby'

# Terminology



# SQL $\rightarrow$ Query Plan


SELECT a FROM R

$\pi_a(R)$




SELECT a FROM R  
WHERE a > 10

$\pi_a(\sigma_{a>10}(R))$



SELECT a  
FROM R JOIN S  
ON R.b = S.b

$\pi_a(\bowtie_b(R,S))$



# Query Evaluation

Naïve execution (operator at a time)

read R

filter  $a > 10$  and write out

read and project a

Cost:  $B + M + M$

SELECT a  
FROM R  
WHERE  $a > 10$

$\pi_a$   
|  
 $\sigma_{a > 10}$   
|  
R

**B** # data pages

**M** # pages matched in  
WHERE clause

Could we do better?

# Query Evaluation

Pipelined exec (tuple/page at a time)

read first page of R, pass to  $\sigma$

filter  $a > 10$  and pass to  $\pi$

project a

(all operators run concurrently)

Cost: B

SELECT a  
FROM R  
WHERE a > 10

$\pi_a$   
|  
 $\sigma_{a>10}$   
|  
R

**B** # data pages

**M** # pages matched in  
WHERE clause


Note: can't pipeline some operators!

e.g., sort, some joins, aggregates


why?

# Predicate Push Down

SELECT a  
FROM R  
WHERE a > 10



(a)



(b)

Which is faster if B+ Tree index: (a) or (b)?

(a)  $\log_F(B) + M$  pages

(b) B pages

It's a Good Idea, especially when we look at Joins

# Nested Loops Join:

```
# outer ⋈1 inner
# outer JOIN inner ON outer.1 = inner.1
for row in outer:
    for irow in inner:
        if row[0] == irow[0]:    # could be any check
            yield (row, irow)
```

## Very flexible

Equality check can be replaced with any condition

Incremental algorithm

Cost:  $M + MN$

Is this the same as a cross product?



# Nested Loops Join

What this means in terms of disk IO

tableA join tableB; tableA is “outer”; tableB is “inner”

M pages in tableA, N pages in tableB, T tuples per page

$$M + T \times M \times N$$

for each tuple  $t$  in tableA, (M pages, TM tuples)

scan through each page  $p_i$  in the inner (N pages)

compare all the tuples in  $p_i$  with  $t$

# Indexed Nested Loops Join

```
for row in outer:  
    for irow in index.get(row[0], []):  
        yield (row, irow)
```

Slightly less flexible

Only supports conditions that the index supports

# Sort Merge Join

Sort outer and inner tables on join key

Cost: 2-3 scans of each table

Merge the tables and compute the join

Cost: 1 scan of each table

## Overall Properties

cost:  $3(M+N)$  to  $4(M+N)$

results are sorted

highly sequential access

(weapon of choice for very large datasets)

# Cost Estimation

`estimate(operator, inputs, stats) → cost`

estimate cost for each operator


depends on input *cardinalities* (# tuples)

discussed earlier in lecture

estimate output size for each operator

need to call `estimate()` on inputs!

use selectivity. assume attributes are independent



Try it in PostgreSQL: `EXPLAIN <query>;`

# Estimate Size of Output

Emp: 1000 Cardinality

Dept: 10 Cardinality

Cost(Emp join Dept)

Naïve

# total records	$1000 * 10$	$= 10,000$
Selectivity of Emp	$1 / 1000$	$= 0.001$
Selectivity of Dept	$1 / 10$	$= 0.1$
Join Selectivity	$1 / \max(1k, 10)$	$= 0.001$
Output Card:	$10,000 * 0.001$	$= 10$

note: selectivity defined wrt cross product size

Note: estimate wrong if this is a key/fk join on emp.did = dept.did: 1000 results

# Transactions

Sequence of actions treated as a single unit

**Atomicity:** Apply all changes or none  
(“atomic” because it is indivisible)

Solves the crash problem

**Isolation:** Illusion that each transaction  
executes sequentially, without concurrency

# Transaction Guarantees

## A

### Atomicity

“all or nothing”: All changes applied, or none are  
users never see in-between transaction state

## C

### Consistency

database always satisfies Integrity Constraints  
Transactions move from valid database to valid database

## I

### Isolation:

from transaction's point of view, it's the only one running

## D

### Durability:

if transaction commits, its effects *must persist*

# Concepts

## Serial schedule

One transaction at a time. no concurrency.

## Equivalent schedule

the database state is the same at end of both schedules

## Serializable schedule (gold standard)

equivalent to a serial schedule



# Why Serializable Schedule? Anomalies

Reading in-between (uncommitted) data

T1: R(A) W(A) R(B) W(B) abort

T2: R(A) W(A) commit

WR conflict or dirty reads

Reading same data gets different values

T1: R(A) R(A) W(A) commit

T2: R(A) W(A) commit

RW conflict or unrepeatable reads

# Why Serializable Schedule? Anomalies

# Stepping on someone else's writes

**T1:**     $W(A)$                        $W(B)$  commit

T2:  $W(A) W(B)$  commit

## WW conflict or lost writes

Notice: all anomalies involve writing to data that is read/written to.

## If we track our writes, maybe can prevent anomalies

# Conflict Serializability

*def: a schedule that is conflict equivalent to a serial schedule*

Meaning: you can swap non-conflicting operations to derive a serial schedule.

$\forall$  conflicting operations  $O_1$  of  $T_1$ ,  $O_2$  of  $T_2$   
 $O_1$  always before  $O_2$  in the schedule or  
 $O_2$  always before  $O_1$  in the schedule

# Conflict Serializability

## Transaction Precedence Graph

Edge  $T_i \rightarrow T_j$  if:

1.  $T_i$  read/write  $A$  before  $T_j$  writes  $A$  or
2.  $T_i$  writes some  $A$  before  $T_j$  reads/writes  $A$

If graph is acyclic (does not contain cycles) then conflict serializable!

# Fine, but what about COMMITing?

T1	R(A) W(A)	R(B) ABORT
T2	R(A) COMMIT	

Not recoverable

Promised T2 everything is OK. IT WAS A LIE.

T1	R(A) W(B) W(A)	ABORT
T2	R(A) W(A)	

Cascading Rollback.


T2 read uncommitted data → T1's abort undos T1's ops & T2's

# Two-Phase Locking (2PL)

Growing phase:      acquire locks

Shrinking phase:    release locks

Guarantees serializable schedules!



Uh Oh, same problem

# Lock-based Concurrency Control

Strict two-phase locking (Strict 2PL)

Growing phase: acquire locks

Shrinking phase: release locks

Release locks after commit/abort



Why? Which problem does it prevent?

T1 R(A) W(B) W(A)

ABORT

T2

R(A) W(A)

Avoids cascading rollbacks!

# Deadlocks

T1	R(A) W(A)	W(B)?
T2	R(B) W(B)	W(A)?

Possible for a cycle of transactions to wait forever

Typical solution: abort txn if waiting too long  
(lock timeout)



# Concept: Sequence of changes

Before any change: write it to a sequential file  
(write ahead of the change: write-ahead log)

On commit/abort: write “commit/abort” record

On recovery: replay complete transactions

Requirements: ordered; no “holes”

# Redo-only log

Rule: To write a modified page to disk:

- All writers must have committed  
(no uncommitted changes that might be aborted)
- All log records on disk

On abort: Need to undo changes

Keep “undo” information in memory

# When do disk writes happen?

On commit: Write log to disk, wait to complete (sequential IO)

After commit: Write modified pages when needed, or when idle

“free” the log? Complicated; not covered

# Disadvantage: Big transactions

What about transactions bigger than memory?

Need to write uncommitted changes to disk

Solution: ARIES algorithm (IBM again; 1992)

Idea: Write both undo and redo records

Can write page if undo records on disk

# Redo and Undo log

Rule: To write a modified page to disk:

- Log undo and redo records on disk

On commit: Write all undo/redo records for txn

Write pages at any time after log records on disk!

# Aries Recovery Algorithm

## 3 phases

Analyze the log to find status of all xacts

Committed or in flight?

Redo xacts that were committed

Now at the same state at the point of the crash

Undo partial (in flight) xacts

Recovery is *extremely* tricky and *must be correct*