

Relational Algebra (and SQL)

Reading

Ramakrishnan Sections 4.1 and 4.2

Helpful Reference

https://en.wikipedia.org/wiki/Relational_algebra

Queries: ask your data a question

Declarative language: ask *what* you want, now *how*

Meaning: Queries are high level, “readable”

Not Turing complete (can’t execute any algorithm)

Supports easy, efficient access to large databases

Domain specific language for data access

SQL: Data Definition Language (DDL)

Data Manipulation Language (DML)

Basic Single Table SELECT

SELECT (*output*) FROM (*input*)
WHERE (*condition*)

SELECT * FROM Students

SELECT name FROM Students

SELECT name FROM Students WHERE age < 21

SELECT name, login FROM Students WHERE gpa >= 3

sid	name	login	age	gpa
1	eugene	ewu	20	2.5
2	luis	gravano	25	3.5
3	martha	martha	32	3.9

Single Table Semantics

conceptual evaluation method:

1. FROM clause: retrieve Students relation
2. WHERE clause: Check conditions, discard tuples that fail
3. SELECT clause: Delete unwanted fields

Real evaluation is *much* more efficient, but must produce the same answers.

Ambiguous names

E.g. Students: (sid, name, ...)

Enrolled: (sid, cid, grade)

Qualified names: Use table name: Students.age

Rename: AS (optional): shortcuts, ambiguity, clarity

```
SELECT S.sid, S.name FROM Students AS S
```

```
SELECT S.sid, S.name FROM Students S
```

Related data: Multiple tables

What does this return?

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid = E.sid AND  
       E.grade = 'A'
```

Enrolled

sid	cid	grade
1	2	A
1	3	B
2	2	A+

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid = E.sid AND  
       E.grade = 'A'
```

Students

sid	name
1	eugene
2	luis

Result

name	cid
eugene	2

Multi-Table Semantics

- Modify the FROM clause evaluation
 - 1. FROM clause: compute *cross-product* of Students and

Enrolled

sid	cid	grade
1	2	A
1	3	B
2	2	A+

Students

sid	name
1	eugene
2	luis

Cross-product

sid	cid	grade	sid	name
1	2	A	1	eugene
1	3	B	1	eugene
2	2	A+	1	eugene
1	2	A	2	luis
1	3	B	2	luis
2	2	A+	2	luis

Multi-Table Semantics

Modify the FROM clause evaluation

1. FROM clause: compute *cross-product* of Students, Enrolled
2. WHERE clause: Check conditions, discard tuples that fail
3. SELECT clause: Delete unwanted fields

Formal Relational Query Languages

Formal basis for real languages e.g., SQL

Relational Algebra

Function of operations applied to relations

Operational: step-by-step execution plans

Relational Calculus

Logical, describes what data users want

(not operational, fully declarative)

Definitions

Relation (for this lecture)

Instance is a set of tuples

Schema defines field names and types (domains)

Students(sid int, name text, major text, gpa int)

Fields: Reference by name: (e.g. major)

Reference by position starting at 1: (e.g. 3)

Names are for humans; positions for “intermediate” results

Definitions

Query is a function over **relation instances**

$$Q(R_1, \dots, R_n) = R_{\text{result}}$$

Schemas of input and output relations are *fixed* and well defined by the query Q.

Relational Algebra Overview

Core 5 operations

PROJECT (π)

SELECT (σ)

UNION (\cup)

SET DIFFERENCE ($-$)

CROSSPRODUCT (\times)

Additional operations

RENAME (ρ)

INTERSECT (\cap)

JOIN (\bowtie)

DIVIDE ($/$)

Instances Used Today: Library

Students, Reservations

RI

sid	rid	day
1	101	10/10
2	102	11/11

Use positional or named field notation

SI

sid	name	gpa	age
1	eugene	4	20
2	barak	3	21
3	trump	2	88

Field names in results inherited from input relations (unless specified)

S2

sid	name	gpa	age
4	aziz	3.2	21
2	barak	3	21
3	trump	2	88
5	rusty	3.5	21

Project

$$\pi_{\langle \text{attr1}, \dots \rangle}(A) = R_{\text{result}}$$

Extract desired fields (subset of columns)

Schema is subset of input schema in the projection list

$\pi_{\langle a, b, c \rangle}(A)$ has output schema (a, b, c) w/ types carried over

Project

S2

sid	name	gpa	age
4	aziz	3.2	21
2	barak	3	21
3	trump	2	88
5	rusty	3.5	21

$\pi_{\text{name,age}}(S2) =$

name	age
aziz	21
barak	21
trump	88
rusty	21

Project

S2

sid	name	gpa	age
4	aziz	3.2	21
2	barak	3	21
3	trump	2	88
5	rusty	3.5	21

$\pi_{\text{age}}(S2) =$

age
21
88

Where did all the rows go?

Real systems typically don't remove duplicates. Why?

Select

$$\sigma_{\langle p \rangle}(A) = R_{\text{result}}$$

Select subset of rows that satisfy condition p

Won't have duplicates in result. Why?

Result schema same as input

Select

S1

sid	name	gpa	age
1	eugene	4	20
2	barak	3	21
3	trump	2	88

$$\sigma_{\text{age} < 30} (S1) =$$

sid	name	gpa	age
1	eugene	4	20
2	barak	3	21

$$\pi_{\text{name}}(\sigma_{\text{age} < 30} (S1)) =$$

name
eugene
barak

Commutativity

$$A + B = B + A$$

$$A * B = B * A$$

$$A + (B * C) = (B * C) + A$$

Associativity:

$$A + (B + C) = (A + B) + C$$

$$A + (B * C) = (A + B) * C$$

Commutativity

$$A + B = B + A$$

$$A * B = B * A$$

$$A + (B * C) = (B * C) + A$$


Associativity:

$$A + (B + C) = (A + B) + C$$


~~$$A + (B * C) = (A + B) * C$$~~


Commutativity

$$\pi_{\text{age}}(\sigma_{\text{age} < 30} (SI))$$

$\sigma_{\text{age} < 30}$ 

sid	name	gpa	age
1	eugene	4	20
2	barak	3	21
3	trump	2	88






sid	name	gpa	age
1	eugene	4	20
2	barak	3	21

Commutativity

$$\pi_{\text{age}}(\sigma_{\text{age} < 30}(SI))$$

$\sigma_{\text{age} < 30}$ 

sid	name	gpa	age
1	eugene	4	20
2	barak	3	21
3	trump	2	88

π_{age} 

sid	name	gpa	age
1	eugene	4	20
2	barak	3	21

=

age
20
21

Commutativity

$$\sigma_{\text{age} < 30}(\pi_{\text{age}}(SI))$$

π_{age} (

sid	name	gpa	age
1	eugene	4	20
2	barak	3	21
3	trump	2	88

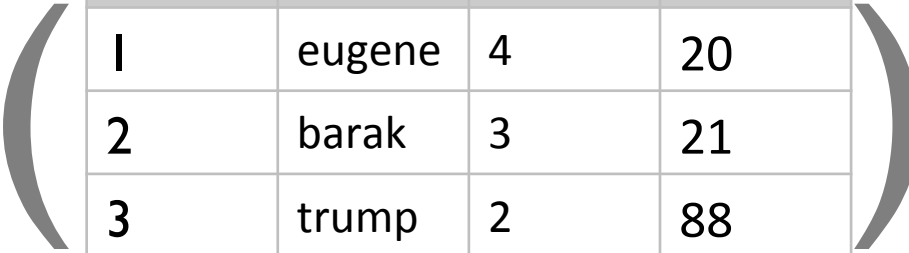
)

}


age
20
21
88


Commutativity

$$\sigma_{\text{age} < 30}(\pi_{\text{age}}(SI))$$


π_{age} 

sid	name	gpa	age
1	eugene	4	20
2	barak	3	21
3	trump	2	88



$\sigma_{\text{age} < 30}$ 

age
20
21
88



=

age
20
21

Commutativity

Does Project and Select commute?

$$\pi_{\text{age}}(\sigma_{\text{age} < 30} (SI)) = \sigma_{\text{age} < 30}(\pi_{\text{age}}(SI))$$

What about

$$\pi_{\text{name}}(\sigma_{\text{age} < 30} (SI))?$$

Commutativity

Does Project and Select commute?

$$\pi_{\text{age}}(\sigma_{\text{age} < 30}(SI)) = \sigma_{\text{age} < 30}(\pi_{\text{age}}(SI))$$

What about

$$\pi_{\text{name}}(\sigma_{\text{age} < 30}(SI)) \neq \sigma_{\text{age} < 30}(\pi_{\text{name}}(SI))$$

Commutativity

Does Project and Select commute?

$$\pi_{\text{age}}(\sigma_{\text{age} < 30}(SI)) = \sigma_{\text{age} < 30}(\pi_{\text{age}}(SI))$$

What about

$$\pi_{\text{name}}(\sigma_{\text{age} < 30}(SI)) \neq \sigma_{\text{age} < 30}(\pi_{\text{name, age}}(SI))$$

Commutativity

Does Project and Select commute?

$$\pi_{\text{age}}(\sigma_{\text{age} < 30}(SI)) = \sigma_{\text{age} < 30}(\pi_{\text{age}}(SI))$$

What about

$$\pi_{\text{name}}(\sigma_{\text{age} < 30}(SI)) = \pi_{\text{name}}(\sigma_{\text{age} < 30}(\pi_{\text{name, age}}(SI)))$$

OK!

Union, Set-Difference

$$A \text{ op } B = R_{\text{result}}$$

A, B must be *union-compatible*

Same number of fields

Field i in each schema have same type

Result Schema borrowed from first arg (A)

Student(sid int, age int) U Class(cid int, max int) = ?

Union, Set-Difference

$$A \text{ op } B = R_{\text{result}}$$

A, B must be *union-compatible*

Same number of fields

Field i in each schema have same type

Result Schema borrowed from first arg (A)

Student(sid int, age int) U Class(cid int, max int) =

R_{result} (sid int, age int)

Union, Intersect, Set-Difference

S1

sid	name	gpa	age
1	eugene	4	20
2	barak	3	21
3	trump	2	88

S2

sid	name	gpa	age
4	aziz	3.2	21
2	barak	3	21
3	trump	2	88
5	rusty	3.5	21

$S1 \cup S2 =$

sid	name	gpa	age
1	eugene	4	20
4	aziz	3.2	21
5	rusty	3.5	21
3	trump	2	88
2	barak	3	21

Union, Intersect, Set-Difference

S1

sid	name	gpa	age
1	eugene	4	20
2	barak	3	21
3	trump	2	88

S2

sid	name	gpa	age
4	aziz	3.2	21
2	barak	3	21
3	trump	2	88
5	rusty	3.5	21

$S1 - S2 =$

sid	name	gpa	age
1	eugene	4	20

Note on Set Difference & Performance

Most operators are monotonic
increasing size of inputs \rightarrow outputs grow
can compute *incrementally*

Set Difference is *not monotonic*:

compute $A - B$

add data to B: $B2 = B \cup X$

compute $A - B2$

could be smaller than $A - B$

Set difference is *blocking*:

For $T - S$, must wait for all S tuples before any results

Intersect: $A \cap B$

Keep tuples in both A and B

Can be defined with Set Difference

$$A \cap B = A - (A - B)$$

Cross-Product

$$A(a_1, \dots, a_n) \times B(a_{n+1}, \dots, a_m) = R_{\text{result}}(a_1, \dots, a_m)$$

Each row of A paired with each row of B

Result schema: Combine A and B's fields, inherit if possible

Conflict: students and reservations have *sid* field

Cross-Product

SI

sid	name	gpa	age
1	eugene	4	20
2	barak	3	21
3	trump	2	88

RI

sid	rid	day
1	101	10/10
2	102	11/11

SI x RI =

(sid)	name	gpa	age	(sid)	rid	day
1	eugene	4	20	1	101	10/10
2	barak	3	21	1	101	10/10
3	trump	2	88	1	101	10/10
1	eugene	4	20	2	102	11/11
2	barak	3	21	2	102	11/11
3	trump	2	88	2	102	11/11

Rename

$\rho(<\text{new_name}>(<\text{mappings}>), Q)$

Explicitly defines/changes field names of schema

$\rho(C(1 \rightarrow \text{sid1}, 5 \rightarrow \text{sid2}), S1 \times R1)$

C =

sid1	name	gpa	age	sid2	rid	day
1	eugene	4	20	1	101	10/10
2	barak	3	21	1	101	10/10
3	trump	2	88	1	101	10/10
1	eugene	4	20	2	102	11/11
2	barak	3	21	2	102	11/11
3	trump	2	88	2	102	11/11