

Structured Query Language  
SQL Es-Que-EI or Sequel

# Announcements

HW1: Returned on Thursday

HW2: Available Thursday

Azure codes available

Next week:

Project I Part 2: Feb 25 (next week Thursday)

HW2: Due March 1st (**Tuesday**)

Midterm: Covers SQL; Will post last Midterm

Review session: Send requests

# Didn't we already talk about SQL?

## Two sublanguages

### **DDL** Data Definition Language

define and modify schema (physical, logical, view)

CREATE TABLE, Integrity Constraints

### **DML** Data Manipulation Language

get and modify data

simple SELECT, INSERT, DELETE

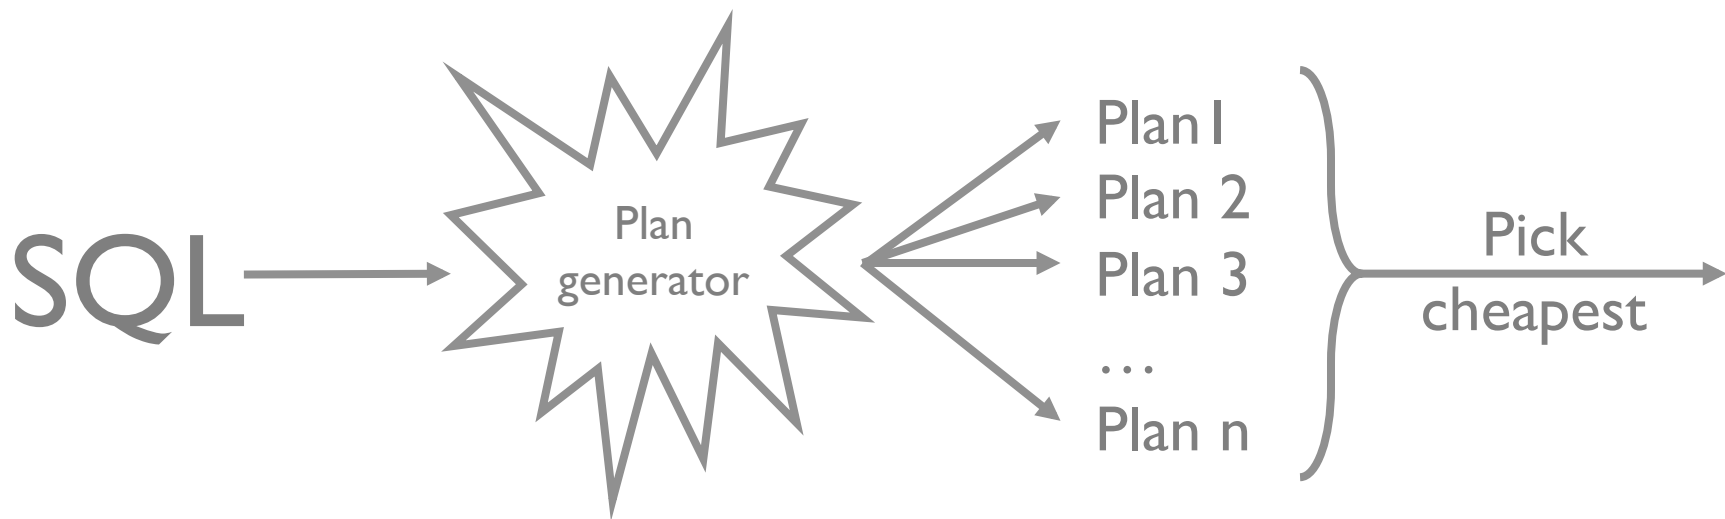
*human-readable* language

# DBMS (tries to) execute efficiently

Key: precise query semantics

Reorder/modify queries while answers stay same

DBMS estimates costs for different evaluation plans



# SQL: Extended Relational Algebra

Multisets rather than sets

Relations can contain duplicates (unless constrained)

Order doesn't matter

NULLs

Aggregates

# Today's Database

Sailors

<u>sid</u>	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Boats

<u>bid</u>	name	color
101	Legacy	red
102	Melon	blue
103	Mars	red

Reserves

<u>sid</u>	<u>bid</u>	day
1	102	9/12
2	102	9/13
2	103	9/14

Is Reserves table correct?

# Today's Database

Sailors

<u>sid</u>	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Boats

<u>bid</u>	name	color
101	Legacy	red
102	Melon	blue
103	Mars	red

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13
2	103	9/14

Is Reserves table correct?  
Day should be part of key

# Today's Database

Sailors

<u>sid</u>	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Boats

<u>bid</u>	name	color
101	Legacy	red
102	Melon	blue
103	Mars	red

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13
2	103	9/14
2	103	9/15

Is Reserves table correct?  
Day should be part of key



# <30 year old sailors

```
SELECT *  
FROM Sailors  
WHERE age < 30
```

<u>sid</u>	name	rating	age
1	Eugene	7	22
3	Ken	8	27

```
SELECT name, age  
FROM Sailors  
WHERE age < 30
```

name	age
Eugene	22
Ken	27

# <30 year old sailors

```
SELECT *  
FROM Sailors  
WHERE age < 30
```

$\sigma_{\text{age} < 30} (\text{Sailors})$

```
SELECT name, age  
FROM Sailors  
WHERE age < 30
```

$\pi_{\text{name, age}} (\sigma_{\text{age} < 30} (\text{Sailors}))$

# Who has reserved boat 102?

Sailors

<u>sid</u>	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

<u>sid</u>	<u>bid</u>	day
1	102	9/12
2	102	9/13
2	103	9/14

# Who has reserved boat 102?

```
SELECT S.name
FROM   Sailors AS S, Reserves AS R
WHERE  S.sid = R.sid AND R.bid = 102
```


Sailors

<u>sid</u>	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13
2	103	9/14

name
Eugene
Luis



# Who reserved boat 102?

```
SELECT S.name  
FROM   Sailors AS S, Reserves AS R  
WHERE  S.sid = R.sid AND R.bid = 102
```

$\pi_{\text{name}} (\sigma_{\text{bid}=102} (\text{Sailors} \bowtie_{\text{sid}} \text{Reserves}))$

(equi-join)

# Who has reserved boat 102?

```
SELECT S.name
FROM   Sailors AS S, Reserves AS R
WHERE  S.sid = R.sid AND R.bid = 102
```

Sailors

<u>sid</u>	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13
2	103	9/14
<b>1</b>	<b>102</b>	<b>9/15</b>

name
Eugene
Luis
Eugene

# DISTINCT: unique rows / set

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13
2	103	9/14

```
SELECT  bid
FROM    Reserves
```

<u>bid</u>
102
102
103

```
SELECT  DISTINCT bid
FROM    Reserves
```

<u>bid</u>
102
103

# Structure of a SQL Query

## DISTINCT

Optional: Remove duplicates (set)

Default: duplicates permitted (multiset)

## target-list

List of expressions over attrs of tables in relation-list

```
SELECT  [DISTINCT] target-list
FROM    relation-list
WHERE   qualification
```

## relation-list

List of relation names

Can define aliases “AS X”

## qualification

Boolean expressions

Combined w/ AND, OR, NOT

attr op const

attr<sub>1</sub> op attr<sub>2</sub>

op is =, <, >, <>, etc




# Semantics

SELECT    [DISTINCT] *target-list*  
FROM       *relation-list*  
WHERE     *qualification*

FROM	compute cross product of relations
WHERE	remove tuples that fail qualifications
SELECT	remove fields not in target-list
DISTINCT	remove duplicate rows

# Conceptual Query Evaluation

**SELECT** [DISTINCT] *target-list*  
**FROM** *relation-list*  
**WHERE** *qualification*  
*GROUP BY* *grouping-list*  
*HAVING* *group-qualification*



Not how actually executed! Above is likely very slow

# Sailors that reserved 1+ boats

```
SELECT  S.sid  
FROM    Sailors AS S, Reserves AS R  
WHERE   S.sid = R.sid
```

Would DISTINCT change anything in this query?

Sailors.sid is a primary key

What if SELECT clause was SELECT S.name?

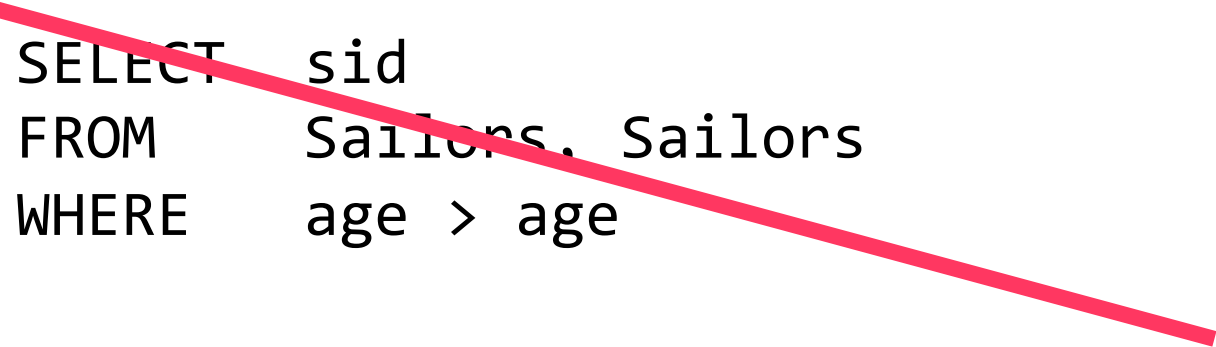
# Sailors that reserved 1+ boats

```
SELECT  DISTINCT S.sid  
FROM    Sailors AS S, Reserves AS R  
WHERE   S.sid = R.sid
```

# Table Alias (AS, Range Variables)

Disambiguate relations

same table used multiple times (self join)



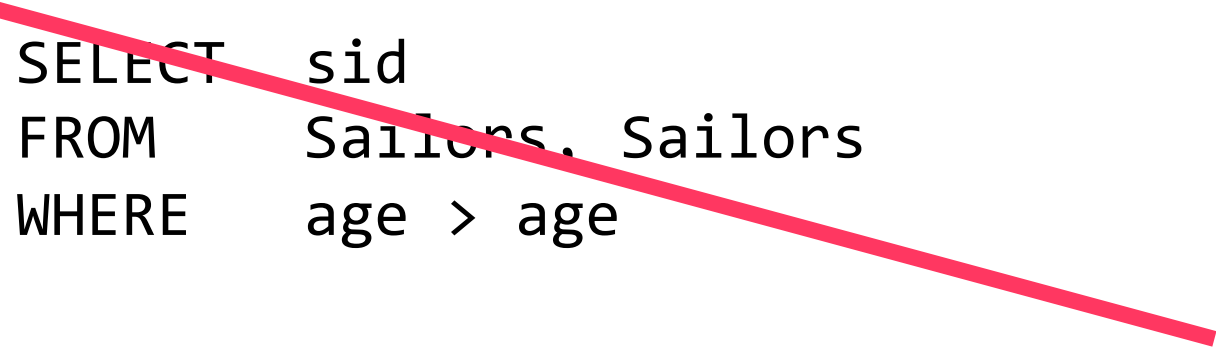
```
SELECT sid  
FROM Sailors, Sailors  
WHERE age > age
```

```
SELECT S1.sid  
FROM Sailors AS S1, Sailors AS S2  
WHERE S1.age > S2.age
```

# Table Alias (AS, Range Variables)

Disambiguate relations

same table used multiple times (self join)



```
SELECT sid  
FROM Sailors, Sailors  
WHERE age > age
```

```
SELECT S1.name, S1.age, S2.name, S2.age  
FROM Sailors AS S1, Sailors AS S2  
WHERE S1.age > S2.age
```

# Expressions (Math)

```
SELECT  S.age, S.age - 5 AS age2, 2*S.age AS age3
FROM    Sailors AS S
WHERE   S.name = 'eugene'
```

```
SELECT  S1.name AS name1, S2.name AS name2
FROM    Sailors AS S1, Sailors AS S2
WHERE   S1.rating*2 = S2.rating - 1
```

# Expressions (Strings)

```
SELECT  S.name
FROM    Sailors AS S
WHERE   S.name LIKE 'e_%'
```

Strings quoted with single quotes: ' (identifiers: double quote)  
If you need an embedded quote: use two: 'this is "quoted"'

'\_' any one character (• in regex)

'%' 0 or more characters of any kind (•\* in regex)

Most DBMSes have rich string manipulation support e.g., regex

PostgreSQL documentation

<http://www.postgresql.org/docs/9.3/static/functions-string.html>



# Expressions (Date/Time)

```
SELECT  R.sid  
FROM    Reserves AS R  
WHERE   now() - R.date < interval '1 day'
```

TIMESTAMP, DATE, TIME types

Values quoted: '2016-02-16', 'Feb-16-2016', '4:05 PM'

now() returns timestamp at start of transaction


DBMSes provide rich time manipulation support

exact support may vary by vendor

Postgresql Documentation

<http://www.postgresql.org/docs/9.3/static/functions-datetime.html>

# Expressions

Constant	1, 'hello', 7.85
Col reference	Sailors.name
Arithmetic	Sailors.sid * 10
Unary operators	NOT
Binary operators	AND, OR, <, =,  , >=
Function	abs(), sqrt(), ...
Casting	1.7::int, '10-12-2015'::date

# UNION, INTERSECT, EXCEPT

Algebra:  $\cup$ ,  $\cap$ ,  $-$

Combine results from two queries:

```
SELECT [query1] UNION SELECT [query2]
```

By default: *distinct results!* (set semantics)

(*operator*) ALL: Keep duplicates: multi-set