

More SQL: Nested Queries

Today's Database

Sailors

<u>sid</u>	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Boats

<u>bid</u>	name	color
101	Legacy	red
102	Melon	blue
103	Mars	red

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13
2	103	9/14
2	103	9/15

Is Reserves table correct?
Day should be part of key

Today's Database

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13
2	103	9/14
2	103	9/15

PRIMARY KEY (sid, bid)

Sailor can only reserve a boat (e.g. 102) **once**

PRIMARY KEY (sid, bid, day)

Boat (e.g. 102) reserved by 2 sailors on same day

Today's Database

PRIMARY KEY (sid, bid, day)

Boat (e.g. 102) reserved by 2 sailors on same day

+ UNIQUE (bid, day)? Works!

PRIMARY KEY (sid, bid, day) + UNIQUE (bid, day) =
PRIMARY KEY(bid, day) + sid NOT NULL

HWI bugs

Missing CHECK constraints

```
Prof(  
    type text,  
    check(text = 'junior' or text = 'senior'),  
)
```

CHECK constraints

Useful for single table constraints

Uncommon in reality, but not unheard of

Remember that they exist for tests

Bank: CHECK(balance > 0)

School: CHECK(
isInstructor OR degree IS NOT NULL)

UNION, INTERSECT, EXCEPT

Algebra: \cup , \cap , $-$

Combine results from two queries:

```
SELECT [query1] UNION SELECT [query2]
```

By default: *distinct results!* (set semantics)

(*operator*) ALL: Keep duplicates: multi-set

sid of Sailors that reserved red or blue boat

```
SELECT  DISTINCT R.sid
FROM    Boats B, Reserves R
WHERE   B.bid = R.bid AND
        (B.color = 'red' OR B.color = 'blue')
```

OR

```
SELECT  R.sid
FROM    Boats B, Reserves R
WHERE   B.bid = R.bid AND B.color = 'red'
UNION
SELECT  R.sid
FROM    Boats B, Reserves R
WHERE   B.bid = R.bid AND B.color = 'blue'
```


sid of Sailors that reserved red or blue boat

```
SELECT    R.sid
FROM      Boats B, Reserves R
WHERE     B.bid = R.bid AND
          (B.color = 'red' OR B.color = 'blue')
```

OR

```
SELECT    R.sid
FROM      Boats B, Reserves R
WHERE     B.bid = R.bid AND B.color = 'red'
UNION ALL
SELECT    R.sid
FROM      Boats B, Reserves R
WHERE     B.bid = R.bid AND B.color = 'blue'
```

sid of Sailors that reserved red and blue boat

```
SELECT    R.sid  
FROM      Boats B, Reserves R  
WHERE     B.bid = R.bid AND  
          (B.color = 'red' AND B.color = 'blue')
```

```
SELECT    R.sid  
FROM      Boats B, Reserves R  
WHERE     B.bid = R.bid AND B.color = 'red'  
INTERSECT  
SELECT    R.sid  
FROM      Boats B, Reserves R  
WHERE     B.bid = R.bid AND B.color = 'blue'
```

sid of Sailors that reserved red and blue boat

Can use self-join instead

```
SELECT    DISTINCT R1.sid
FROM      Boats B1, Reserves R1
WHERE
          B1.bid = R1.bid AND

          B1.color = 'red'
```

sid of Sailors that reserved red and blue boat

Can use self-join instead

```
SELECT  DISTINCT R1.sid
FROM    Boats B1, Reserves R1, Boats B2, Reserves R2
WHERE   B1.bid = R1.bid AND

        B1.color = 'red'
```

sid of Sailors that reserved red and blue boat

Can use self-join instead

```
SELECT  DISTINCT R1.sid
FROM    Boats B1, Reserves R1, Boats B2, Reserves R2
WHERE   B1.bid = R1.bid AND
        B2.bid = R2.bid AND
        B1.color = 'red' AND B2.color = 'blue'
```

sid of Sailors that reserved red and blue boat

Can use self-join instead

```
SELECT    DISTINCT R1.sid
FROM      Boats B1, Reserves R1, Boats B2, Reserves R2
WHERE     R1.sid = R2.sid AND
          B1.bid = R1.bid AND
          B2.bid = R2.bid AND
          B1.color = 'red' AND B2.color = 'blue'
```

sids of sailors that haven't reserved a boat

```
SELECT    S.sid  
FROM      Sailors S
```

EXCEPT

```
SELECT    S.sid  
FROM      Sailors S, Reserves R  
WHERE     S.sid = R.sid
```

Nested Queries

```
SELECT  S.sid
FROM    Sailors S
WHERE   S.sid IN (SELECT  R.sid
                  FROM    Reserves R
                  WHERE    R.bid = 101)
```

Many clauses can contain SQL queries
WHERE, FROM, HAVING, SELECT

Conceptual model:

- for each Sailors tuple
- run the subquery and evaluate qualification

Nested Query vs Join

```
SELECT  S.sid
FROM    Sailors S
WHERE   S.sid IN (SELECT  R.sid
                  FROM    Reserves R
                  WHERE   R.bid = 101)
```

```
SELECT  S.sid
FROM    Sailors S, Reserves R
WHERE   S.sid = R.sid AND R.bid = 101
```

What if a student reserved a boat more than once?

Nested: No duplicates

Join: Duplicates

SET Comparison Operators

$x \text{ IN } r$: True if value x appears in r

EXISTS r : True if relation r is not empty (NOT EXISTS)

$x \text{ (operator) ANY } r$: True if $x \text{ (operator)}$ is true for any row in r

E.g. $x \text{ IN } r$ is equivalent to $x = \text{ANY } r$

$x \text{ (operator) ALL } r$: True if $x \text{ (operator)}$ is true for all rows in r

E.g. $x \text{ NOT IN } r$ is equivalent to $x \neq \text{ALL } r$

Reference outer table in nested query

```
SELECT  S.sid
FROM    Sailors S
WHERE   EXISTS (SELECT  *
                  FROM    Reserves R
                  WHERE   R.bid = 101 AND
                        S.sid = R.sid)
```

Outer table referenced in nested query

Conceptual model:

- for each Sailors tuple
- run the subquery and evaluate qualification

Sailors whose rating is greater than
any sailor named “Bobby”

```
SELECT S1.name
FROM   Sailors S1
WHERE  S1.rating > ANY (SELECT  S2.rating
                        FROM    Sailors S2
                        WHERE    S2.name = 'Bobby')
```

How are these different?

```
SELECT S1.name
FROM   Sailors S1
WHERE  S1.rating > ANY (SELECT S2.rating
                        FROM   Sailors S2
                        WHERE  S2.name = 'Bobby')
```

```
SELECT S1.name
FROM   Sailors S1
WHERE  S1.rating > ALL (SELECT S2.rating
                        FROM   Sailors S2
                        WHERE  S2.name = 'Bobby')
```

Rewrite INTERSECT using IN

```
SELECT S.sid  
FROM   Sailors S  
WHERE  S.rating > 2
```

INTERSECT

```
SELECT R.sid  
FROM   Reserves R
```

```
SELECT S.sid  
FROM   Sailors S  
WHERE  S.rating > 2 AND  
       S.sid IN (  
    SELECT R.sid  
    FROM   Reserves R  
       )
```

Similar trick for EXCEPT → NOT IN

What if want *names* instead of sids?

Names are not unique!

Name of sailors that reserved all boats

Hint: All is hard: have “EXISTS” not “FOR ALL”

What about double negation?

reserved all boats \equiv no boat w/out reservation

Can we find boats not reserved by sailor x?

Use that to find sailors who do not have any unreserved boats!

Q1: boats not reserved by Sailor 1

Sailors

<u>sid</u>	name	rating	age
1	Eugene	7	22
2	Luis	2	39
3	Ken	8	27

Boats

<u>bid</u>	name	color
101	Legacy	red
102	Melon	blue
103	Mars	red

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13
2	103	9/14
2	101	9/15

Hint: boats reserved by Sailor 1?

Hint: Use a nested query

Want: sailors who
reserved all boats

Boats reserved by Sailor 1

```
SELECT  DISTINCT r.bid  
FROM    Reserves r  
WHERE   r.sid = 1;
```

Boats not reserved by Sailor 1

```
SELECT  b.bid  
FROM    Boats b  
WHERE   b.bid NOT IN (
```

```
    SELECT  r.bid  
    FROM    Reserves r  
    WHERE   r.sid = 1
```

```
);
```

Q2:

All sailors with
unreserved boats

Want: sailors who
reserved all boats

All sailors with unreserved boats

```
SELECT  s.sid, s.name  
FROM    Sailors s  
WHERE   EXISTS (
```

```
    SELECT  b.bid  
    FROM    Boats b  
    WHERE   b.bid NOT IN (
```

```
        SELECT  r.bid  
        FROM    Reserves r  
        WHERE   r.sid = s.sid
```

```
    )  
);
```

Q:
sailors who have
reserved all boats

All sailors who reserved all boats

```
SELECT    s.sid, s.name
FROM      Sailors s
WHERE     NOT EXISTS (

    SELECT    b.bid
    FROM      Boats b
    WHERE     b.bid NOT IN (

        SELECT    r.bid
        FROM      Reserves r
        WHERE     r.sid = s.sid
    )
);
```

Sailors that reserved all boats (Division)

Hint: double negation

reserved all boats $\equiv \nexists \text{ boat } \nexists \text{ reservation}$

```
SELECT S.name  
FROM   Sailors S  
WHERE  NOT EXISTS (
```

Sailors S where there is not

Sailors that reserved all boats (Division)

Hint: double negation

reserved all boats == \nexists boat \nexists reservation

```
SELECT S.name
FROM   Sailors S
WHERE  NOT EXISTS (SELECT B.bid
                   FROM   Boats B
                   WHERE  NOT EXISTS (
```

Sailors S where there is not

Any boat where there is not

Sailors that reserved all boats (Division)

Hint: double negation

reserved all boats == \nexists boat \nexists reservation

```
SELECT S.name
FROM   Sailors S
WHERE  NOT EXISTS (SELECT B.bid
                   FROM   Boats B
                   WHERE  NOT EXISTS (SELECT R.bid
                                     FROM   Reserves R
                                     WHERE  R.sid = S.sid
                                     AND R.bid = B.bid ))
```

Sailors S where there is not

Any boat where there is not

A reservation by S