

L2I

Query Execution & Optimization Continued

SQL \rightarrow Query Plan


SELECT a FROM R

$\pi_a(R)$




SELECT a FROM R
WHERE a > 10

$\pi_a(\sigma_{a>10}(R))$



SELECT a
FROM R JOIN S
ON R.b = S.b

$\pi_a(\bowtie_b(R, S))$



Query Evaluation

Push vs Pull?

Push

Operators are input-driven

As operator (say reading input table) gets data, push it to parent operator.

Pull

Operators are demand-driven

If parent says “give me next result”, then do the work

Are cursors push or pull?

Query Evaluation

Naïve execution (operator at a time)


read R

filter $a > 10$ and write out

read and project a

Cost: $B + M + M$

SELECT a
FROM R
WHERE $a > 10$



B # data pages

M # pages matched in
WHERE clause

Could we do better?

Query Evaluation

Pipelined exec (tuple/page at a time)

read first page of R, pass to σ

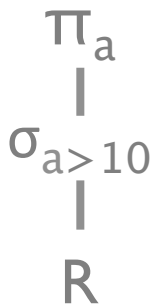
filter $a > 10$ and pass to π

project a

(all operators run concurrently)

Cost: B

SELECT a
FROM R
WHERE a > 10



B # data pages

M # pages matched in
WHERE clause

Note: can't pipeline some operators!

e.g., sort, some joins, aggregates

why?

Query Evaluation

What if R is indexed?

Hash index

Not appropriate

B+Tree index

use $a > 10$ to find initial data page

scan leaf data pages

Cost: $\log_F B + M$

SELECT a
FROM R
WHERE a > 10

π_a
|
 $\sigma_{a > 10}$
|
R

B # data pages

M # pages matched in
WHERE clause

Access Paths

Choice of how to access input data is called the
Access Path

file scan or

index + matching condition (e.g., $a > 10$)

Access Paths

Sequential Scan

doesn't accept any matching conditions

Hash index search key $\langle a, b, c \rangle$

accepts conjunction of equality conditions on *all* search keys

e.g., $a = 1$ and $b = 5$ and $c = 5$

will $(a = 1 \text{ and } b = 5)$ work? why?

Tree index search key $\langle a, b, c \rangle$

accepts conjunction of terms of *prefix* of search keys

typically best with equality on all but last column

e.g., $a = 1$ and $b = 5$ and $c < 5$

will $(a = 1 \text{ and } b > 5)$ work?

will $(a > 1 \text{ and } c > 9)$ work?

How to pick Access Paths?

Selectivity

ratio of # outputs satisfying predicates vs # inputs

0.01 means 1 output tuple for every 100 input tuples


Assume

attribute selectivity is independent

if $\text{selectivity}(a=1) = 0.1$, $\text{selectivity}(b>3) = 0.6$

$\text{selectivity}(a=1 \text{ and } b>3) = 0.1 * 0.6 = 0.06$

High level index structure



What is a data entry?

actual data record

<search key value, rid>

<search key value, rid_list>

How to pick Access Paths?

Hash index on $\langle a, b, c \rangle$

$a = 1, b = 1, c = 1$ how to estimate selectivity?

1. pre-compute attribute statistics by scanning data
e.g., a has 100 values, b has 200 values, c has 1 value
selectivity = $1 / (100 * 200 * 1)$
2. How many distinct values does hash index have?
e.g., 1000 distinct values in hash index
3. make a number up
“default estimate” is the fancy term

System Catalog Keeps Statistics

System R

NCARD	"relation cardinality" # tuples in relation
TCARD	# pages relation occupies
ICARD	# keys (distinct values) in index
NINDX	pages occupied by index
min and max keys in indexes	

Statistics were expensive in 1979!

Catalog stored in relations too!

What Optimization Options Do We Have?

Access Path ✓

Predicate push-down


Join implementation

Join ordering


In general, depends on operator implementations. So let's take a look

Predicate Push Down

SELECT a
FROM R
WHERE a > 10



(a)



(b)

Which is faster if B+ Tree index: (a) or (b)?

(a) $\log_F(B) + M$ pages

(b) B pages

It's a Good Idea, especially when we look at Joins

Projection with DISTINCT clause

need to deduplicate e.g., π_{rating} Sailors

Two basic approaches

Sort: fundamental database operation

sort on rating, remove dups on scan of sorted data

Hash:

partition into N buckets

remove duplicates on insert

Index on projected fields

scan the index pages, avoid reading data

The Join

Core database operation

join of 10s of tables common in enterprise apps

Join algorithms is a large area of research

e.g., distributed, temporal, geographic, multi-dim, range, sensors, graphs, etc

Discuss three basic joins

nested loops, indexed nested loops, hash join

Best join implementation depends on the query, the data, the indices, hardware, etc

Nested Loops Join:

```
# outer ⋈1 inner
# outer JOIN inner ON outer.1 = inner.1
for row in outer:
    for irow in inner:
        if row[0] == irow[0]:    # could be any check
            yield (row, irow)
```

Very flexible

Equality check can be replaced with any condition

Incremental algorithm

Cost: $M + MN$

Is this the same as a cross product?

Nested Loops Join

What this means in terms of disk IO

tableA join tableB; tableA is “outer”; tableB is “inner”

M pages in tableA, N pages in tableB, T tuples per page

$$M + T \times M \times N$$

for each tuple t in tableA, (M pages, TM tuples)

scan through each page p_i in the inner (N pages)

compare all the tuples in p_i with t

Nested Loops Join: Order?

Does order matter?

$$M + T \times M \times N$$

$$N + T \times N \times M$$

Scan “outer” once; Scan “inner” multiple times:

If inner is small IO cost is $M + N$!

Indexed Nested Loops Join

```
for row in outer:  
    for irow in index.get(row[0], []):  
        yield (row, irow)
```

Slightly less flexible

Only supports conditions that the index supports

Indexed Nested Loops Join

What this means in terms of disk IO

outer join inner on sid

M pages in outer, N pages in inner, T tuples/page

inner has primary key index on sid

Cost of looking up in index is C_i

predicate on outer has 5% selectivity

$$M + T \times M \times 0.05 \times C_i$$

for each tuple t in the outer: (M pages, TM tuples)

if predicate(t): (5% of tuples satisfy pred)

lookup_in_index(t.sid) (C_i disk IO)