

Transactions, Concurrency, Recovery

Announcements

Next Tuesday: Guest Lecture

Dr. Zahi Karam, Data Scientist, Bluecore

Homework 4: Due Thursday 8:40 AM

Project 2: Due next Friday April 22nd

Homework 3, Previous finals on web site

Concepts

Serial schedule

One transaction at a time. no concurrency.

Equivalent schedule

the database state is the same at end of both schedules

Serializable schedule (gold standard)

equivalent to a serial schedule

SQL → R/W Operations

```
UPDATE    accounts  
SET       bal = bal + 1000  
WHERE     bal > 1M
```

Read all balances for every tuple

Update those with balances > 1000

Does the access method matter?

Why Serializable Schedule? Anomalies

Reading in-between (uncommitted) data

T1: R(A) W(A) R(B) W(B) abort

T2: R(A) W(A) commit

WR conflict or dirty reads

Reading same data gets different values

T1: R(A) R(A) W(A) commit

T2: R(A) W(A) commit

RW conflict or unrepeatable reads

Why Serializable Schedule? Anomalies

Stepping on someone else's writes

T1: $W(A)$ $W(B)$ commit

T2: $W(A) \ W(B)$ commit

WW conflict or lost writes

Notice: all anomalies involve writing to data that is read/written to.

If we track our writes, maybe can prevent anomalies

Conflict Serializability

What is a conflict?

For 2 operations, if run in different order, get different results

Conflict?	R	W
R	NO	YES
W	YES	YES

Conflict Serializability

def: a schedule that is conflict equivalent to a serial schedule

Meaning: you can swap non-conflicting operations to derive a serial schedule.

\forall conflicting operations O_1 of T_1 , O_2 of T_2
 O_1 always before O_2 in the schedule or
 O_2 always before O_1 in the schedule

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	w	x	y	z
T2:	R(A)	W(A)	R(B)	W(B)

Logical

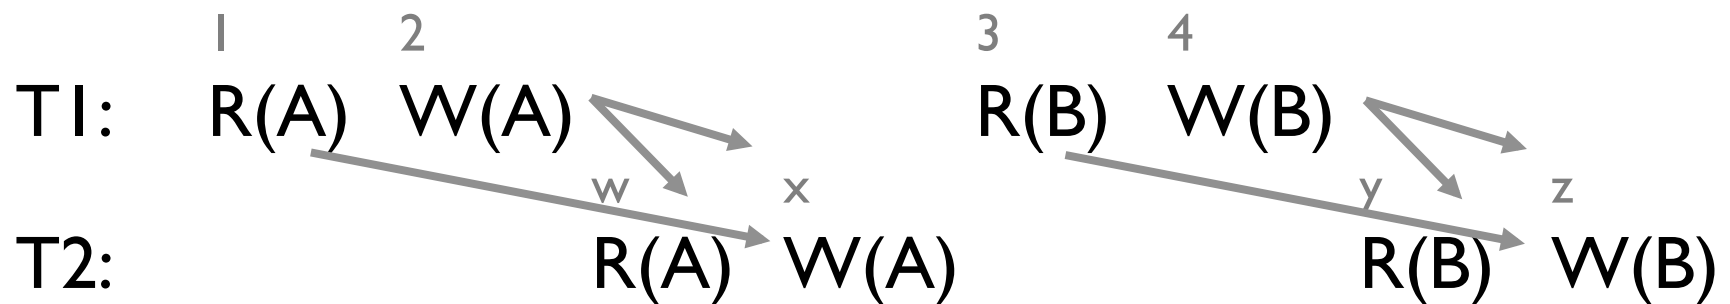
Conflicts

1x, 2w, 2x, 3z, 4y, 4z

Logical

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	w	x	y	z
T2:	R(A)	W(A)	R(B)	W(B)

Serializable




Logical

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	w	x	y	z
T2:	R(A)	W(A)	R(B)	W(B)

Not Serializable

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	w	x	y	z
T2:	R(A)	W(A)	R(B)	W(B)



Conflict Serializability

Transaction Precedence Graph

Edge $T_i \rightarrow T_j$ if:

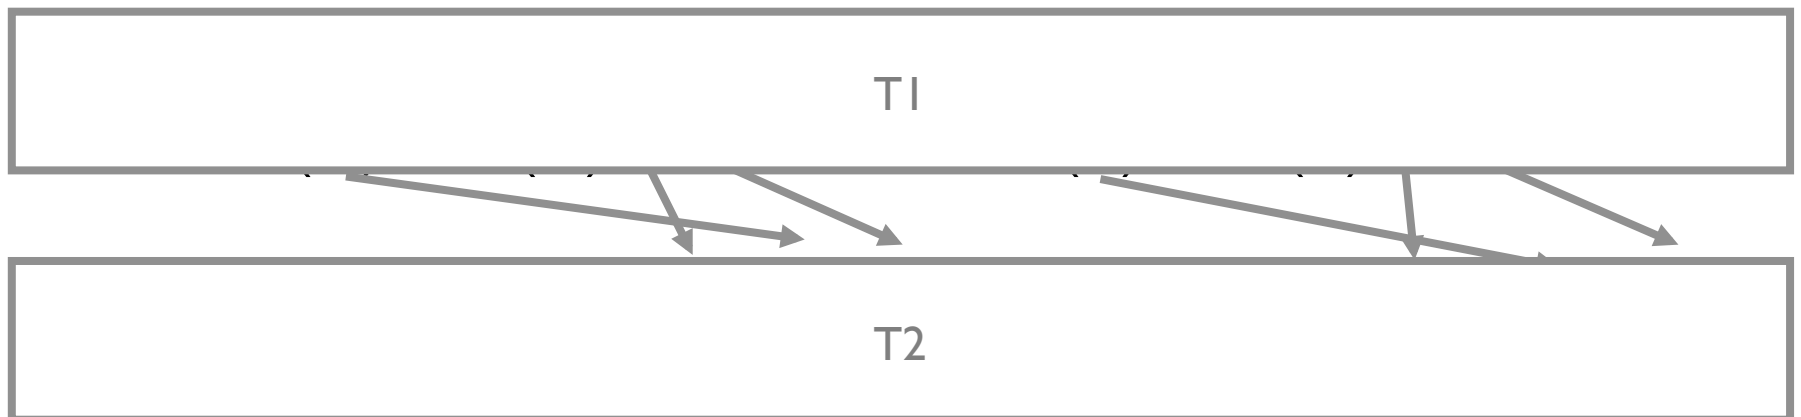
1. T_i read/write A before T_j writes A or
2. T_i writes some A before T_j reads/writes A

If graph is acyclic (does not contain cycles) then conflict serializable!

Logical

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	5	6	7	8
T2:	R(A)	W(A)	R(B)	W(B)

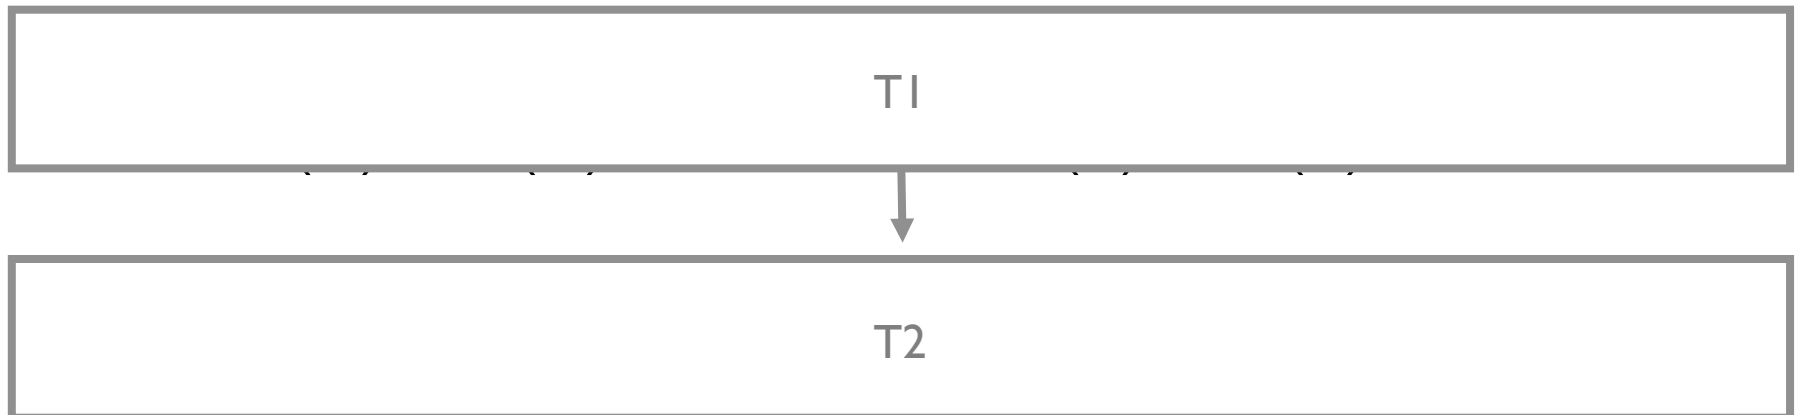
Serializable



Logical

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	5	6	7	8
T2:	R(A)	W(A)	R(B)	W(B)

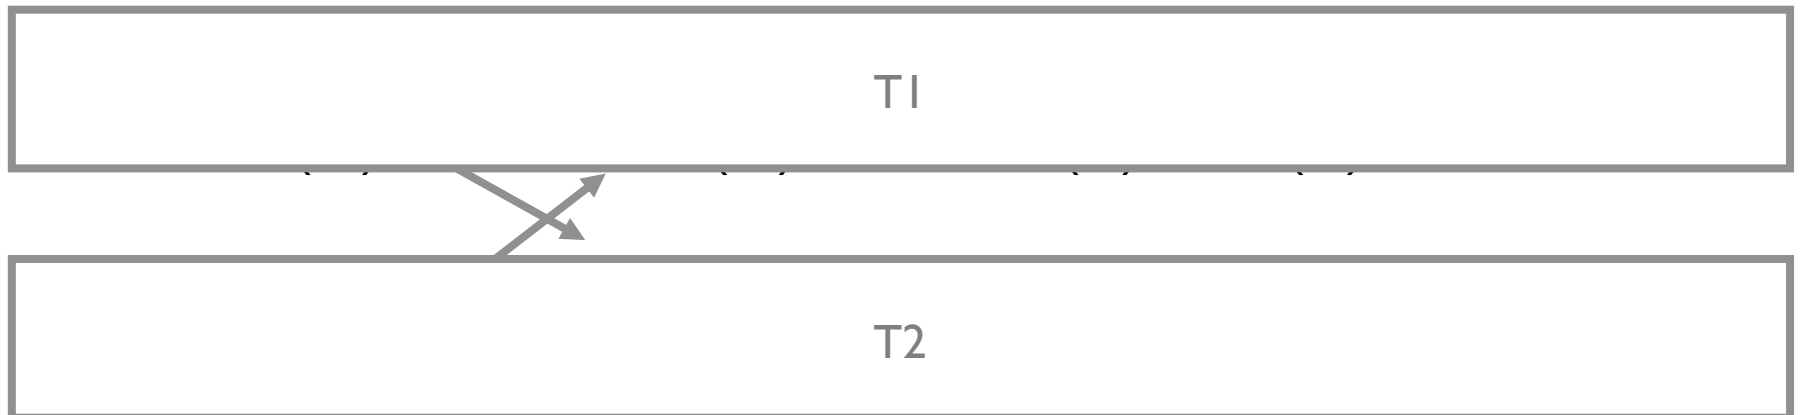
Serializable



Logical

	1	2	3	4
T1:	R(A)	W(A)	R(B)	W(B)
	5	6	7	8
T2:	R(A)	W(A)	R(B)	W(B)

Not Serializable



Fine, but what about COMMITing?

T1	R(A) W(A)	R(B) ABORT
T2	R(A) COMMIT	

Not recoverable

Promised T2 everything is OK. IT WAS A LIE.

T1	R(A) W(B) W(A)	ABORT
T2	R(A) W(A)	

Cascading Rollback.

T2 read uncommitted data → T1's abort undoes T1's ops & T2's

Lock-based Concurrency Control

Lock the object before reading or writing

Read: Other readers permitted, no writers (shared = S)

Write: Must be only transaction accessing (exclusive = X)

		T1	
		S	X
T2	Other Allowed?		
	S	Y	N
	X	N	N

Aside: What is a lock?

An abstraction

Hide complex implementation behind a simple to understand and easy to use interface

Aside: How do locks work?

Single CPU era: Operating system implementation

Process A: Give me access to X

OS: Okay! No one is using it

Process B: Give me access to X

OS: I'll stop running you until A is done

Aside: How do locks work?

Multi CPU era: Special CPU instructions

CPU A: I'm going to write X

CPU B: I'm going to write X

Some hardware somewhere: CPU B wins!

Aside: How do locks work?

How does the hardware work?

No clue

Ancient history example:

CPU A asserts “A exclusive” wire and waits

CPU B asserts “B exclusive” wire and waits

CPU B sees A’s signal and stops; waits and retries

Abstractions

Transactions: Multiple reads/writes

Locks: Exclusive access; wait if not available

CPU instructions: One writer wins

Hardware: Wire delays? Special circuits?

Abstractions

Transactions: Multiple reads/writes

Locks: Exclusive access; wait if not available

CPU instructions: One writer wins

Hardware: Wire delays? Special circuits?

Physics???: Speed of electrons through copper?

Abstractions

Transactions: Multiple reads/writes

Locks: Exclusive access; wait if not available

CPU instructions: One writer wins

Hardware: Wire delays? Special circuits?

Physics???: Speed of electrons through copper?

**Build complex systems out of simple
interfaces**

Lock-based Concurrency Control

Lock the object before reading or writing

Read: Other readers permitted, no writers (shared = S)

Write: Must be only transaction accessing (exclusive = X)

		T1	
		S	X
T2	Other Allowed?		
	S	Y	N
	X	N	N

Can this schedule happen?

T1	R(A)	W(A)		R(B)	ABORT
T2			R(A)	COMMIT	

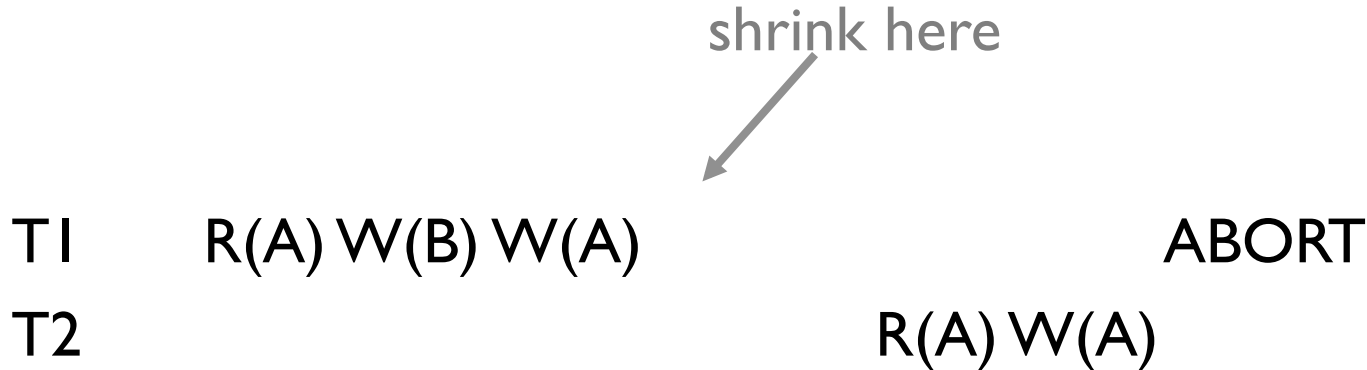
OOPS! Maybe

Two-Phase Locking (2PL)

Growing phase: acquire locks

Shrinking phase: release locks

Guarantees serializable schedules!



Uh Oh, same problem

Lock-based Concurrency Control

Strict two-phase locking (Strict 2PL)

Growing phase: acquire locks

Shrinking phase: release locks

Release locks after commit/abort



Why? Which problem does it prevent?

T1	R(A) W(B) W(A)	ABORT
T2	R(A) W(A)	

Avoids cascading rollbacks!

Deadlocks

T1	R(A) W(A)	W(B)?
T2	R(B) W(B)	W(A)?

Possible for a cycle of transactions to wait forever

Typical solution: abort txn if waiting too long
(lock timeout)

Review

Issues

WR: dirty reads

RW: unrepeatable reads

WW: lost writes

Schedules

Equivalence

Serial

Serializable

Serializability

Conflict serializability

how to detect

Conflict Serializable Issues

Not recoverable

Cascading Rollback

Strict 2 phase locking

Deadlocks

Transactions in SQL

Default: “Autocommit”: Statement = Transaction

BEGIN: Start transaction

COMMIT: Commit transaction

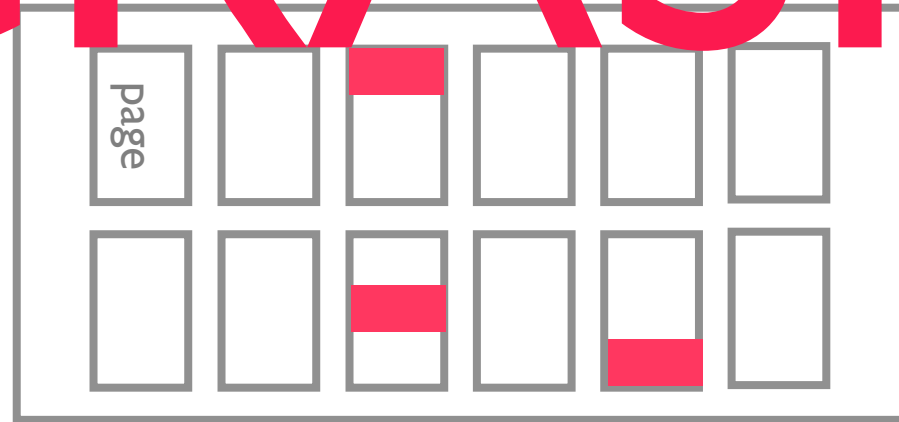
ROLLBACK: Abort transaction

Default “isolation” may be less than serializable

CRASH

Normal Execution

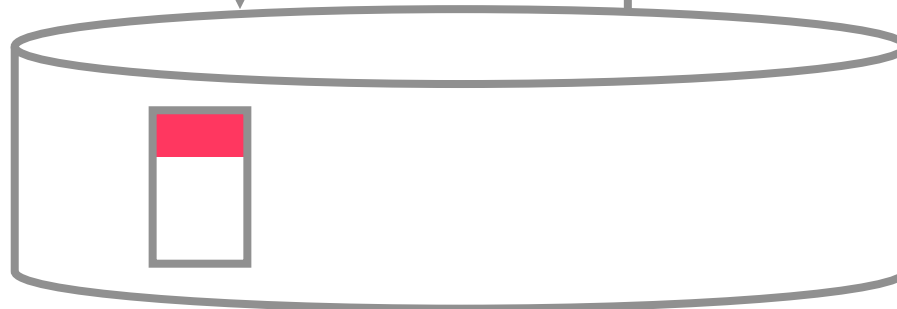
RAM



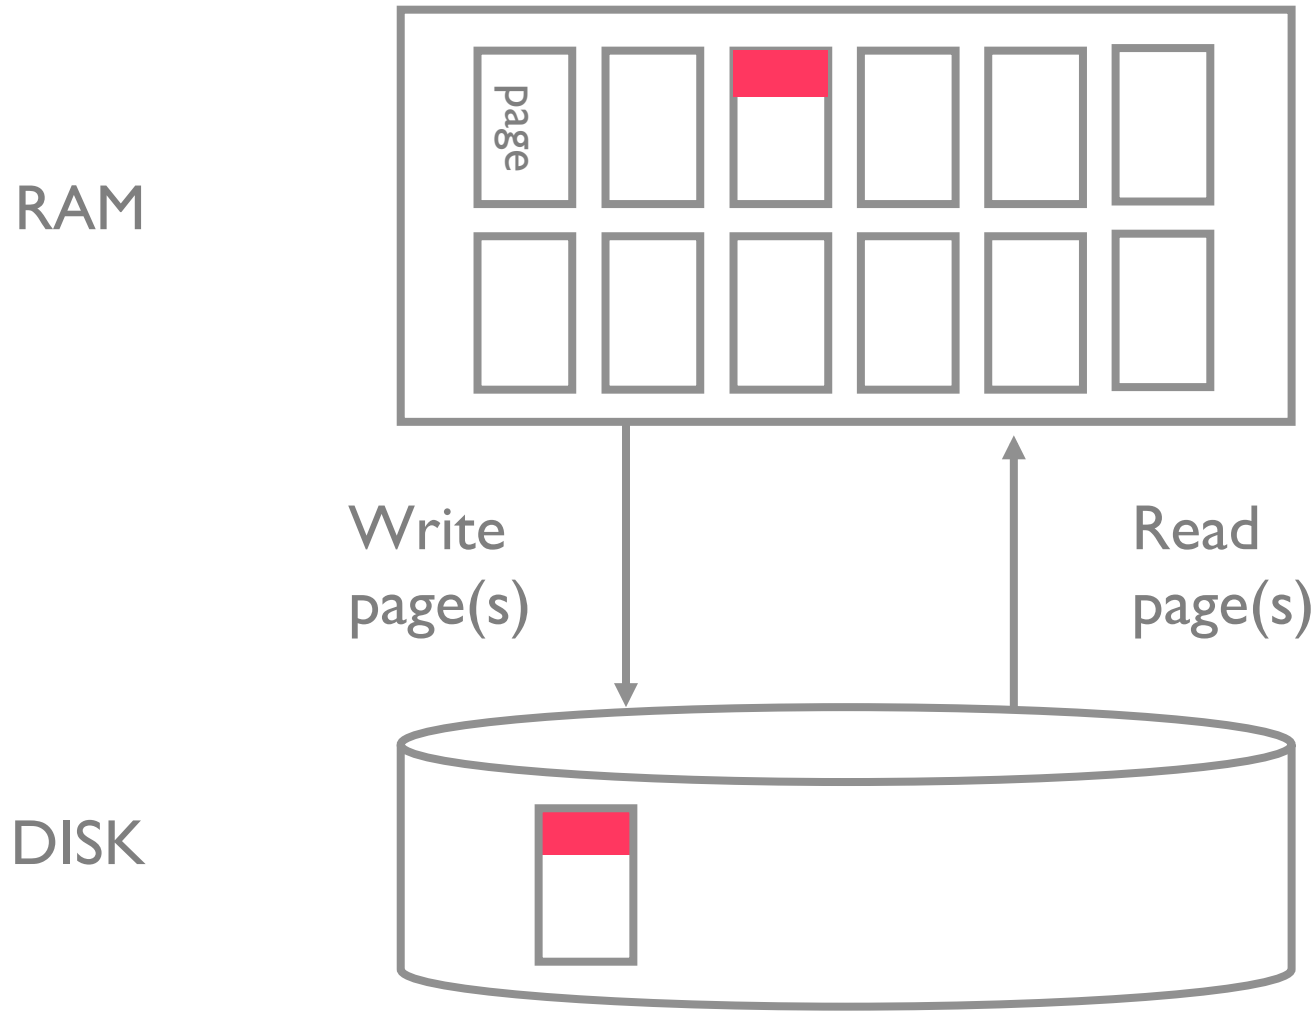
Write
page(s)

Read
page(s)

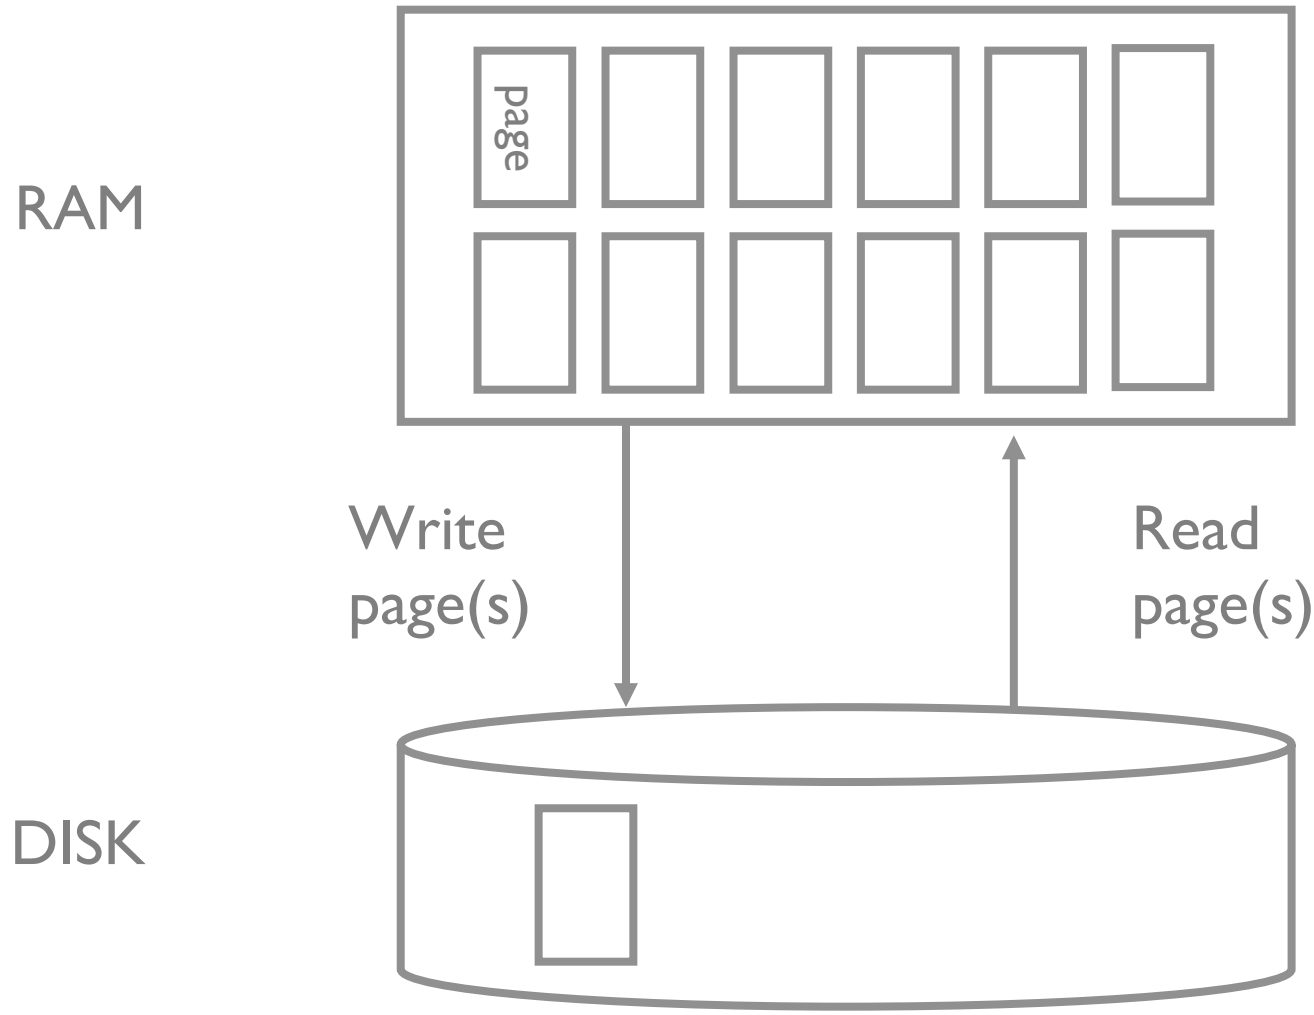
DISK



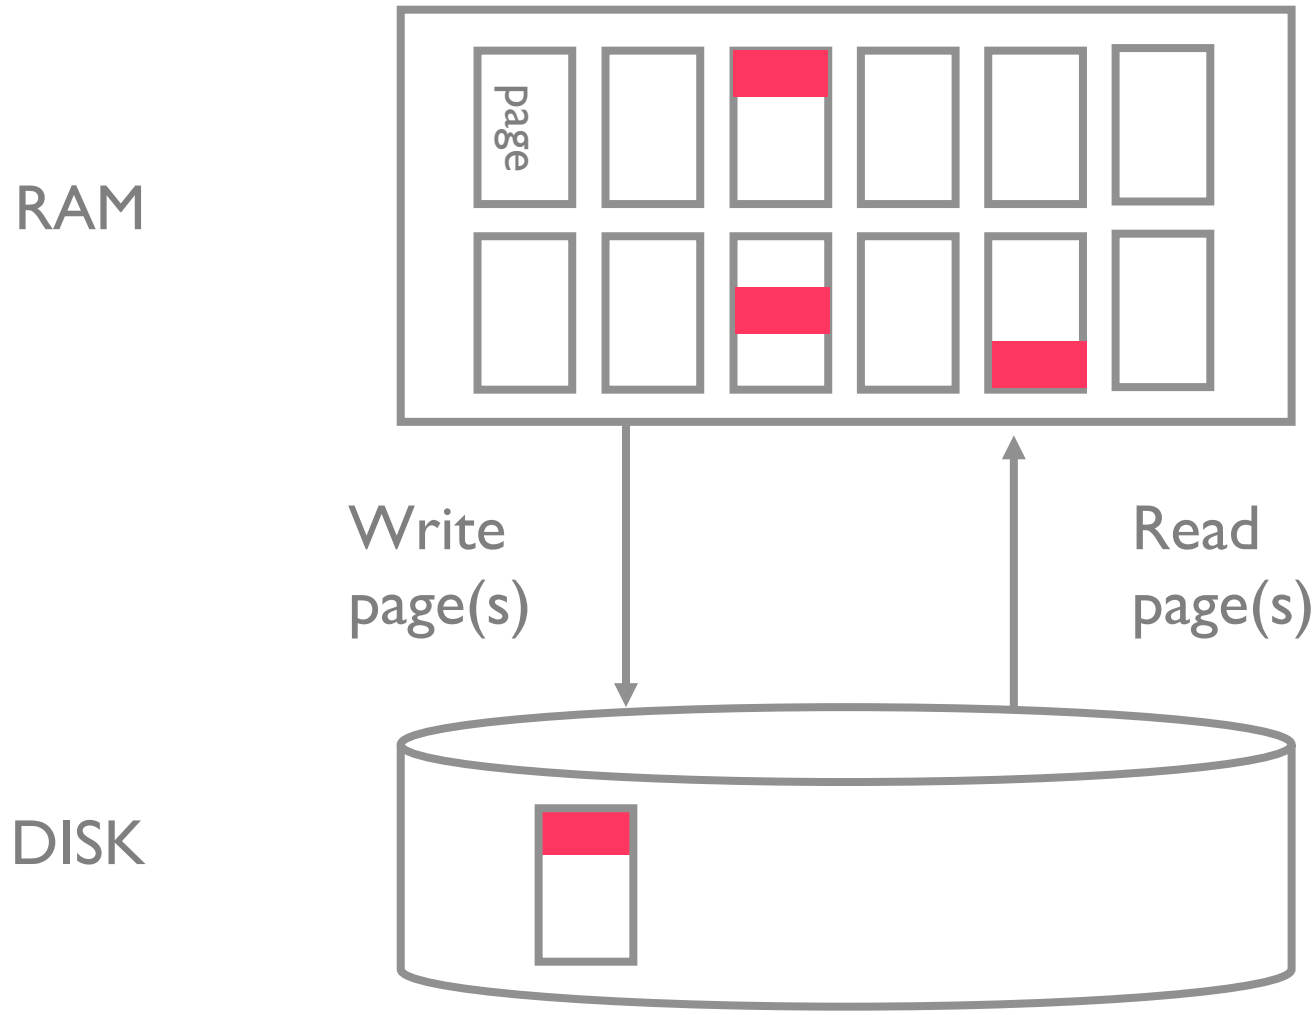
After a Crash



If DB did not say “OK, committed”



If T1 Committed and DB said “OK”



Recovery

Two properties: **Durability**, Atomicity

Assumptions for this class:

Disk is safe. Memory is not.

Running strict-2PL

Need to account for

when pages are modified

when pages are flushed to disk